

Project 2 - SF2957 Statistical Machine Learning

Due December 9, 2025, 12:00

1 Reinforcement learning for blackjack

The purpose of this project is to use reinforcement learning to train an agent to play blackjack. Two different representations of the state space will be considered: one that is based on representing the state as the set of cards and the dealer's card sum, this will be referred to as the “extended state space”, and one that represents the state as the agent's and dealer's card sums. The card sum representation is implemented in some standard approaches to blackjack, but is somewhat flawed (as will be discussed later). Both “on-policy Monte Carlo learning” and “Q-learning” will be considered.

The software environment is already rather complete, so you will only have to implement the core of the learning algorithms and evaluate the results.

2 Blackjack

In this project, the game of blackjack is set up in the following manner.

- One agent plays against the dealer, with the agent staking one unit on each hand/round.
- In each non-terminal state, two actions are possible: ask for another card, or stay.
- The cards 2–10 count as their numerical value, court cards count as 10, and aces count as either 1 or 11 depending on whichever is best. If an ace can be counted as 11 without the agent going bust (card sum exceeds 21), it is known as a *usable* ace; the same goes for the dealer.

The goal of the agent is to beat the dealer in one of the following ways.

- Obtain 21 points on the first two cards, known as a *blackjack*, without a dealer blackjack. Reward: 1.5.
- Reach a final score higher than the dealer without exceeding 21. Reward: 1.

- Dealer gets points exceeding 21, and player does not. Reward: 1.

The dealer always plays according to the same strategy: draw cards until it has a card sum greater than or equal to 17. The game starts with the dealer giving the player two cards (visible) and himself two cards (one visible, one hidden). The agent is then allowed to ask for more cards until she decides to stay, after which the dealer follows his strategy until done. If the player's points exceed 21, her stake is lost regardless of the dealer's outcome (this is where the house edge comes from); if the player's points equal the dealer's points, this is a *push*, and the stake is returned to the player; otherwise, payout is made according to the above rules. Standard practice is that the cards are drawn from 6–10 decks.

2.1 The state space of blackjack

In blackjack, the color and suit of the cards do not matter, only their numerical values. The numerical values will be used as the cards' identifier, with aces equal to 1. The state space may be represented by the agent's hand, a vector (s_1^a, \dots, s_{10}^a) where s_i^a is the number of cards of value i with card sum $\sum_{i=1}^{10} i s_i^a \leq 31$, and the dealer's hand (s_1^d, \dots, s_{10}^d) , where $\sum_{i=1}^{10} i s_i^d \leq 26$. With this state representation, it is always possible to determine if a state is terminal; it is sufficient to check if the agent or dealer is bust, or if the dealer's card sum exceeds 17. Moreover, the hand $(1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1)$ is a blackjack.

By the rules of the game, the agent's policy needs only to take into account the dealer's first card, and the state space can be reduced to

$$\mathcal{S}_1 = \{(s_1, \dots, s_{10}, S_d) : \sum_{i=1}^{10} i s_i \leq 31, S_i \geq 0 \text{ for } i = 1, \dots, 10, \text{ and } S_d \leq 26\},$$

where S_d is the dealer's card sum. The rewards R_1, R_2, \dots follows the above payout rules, with $R_t = 0$ if S_t is not terminal.

An alternative state space, sometimes encountered for blackjack environments, is to further reduce to

$$\mathcal{S}_2 = \{(S_a, u, S_d) : 1 \leq S \leq 31, u \in \{0, 1\}, 1 \leq S_d \leq 26\},$$

where S_a and S_d are the card sums of the agent and the dealer, respectively, and u indicates if the player is holding a usable ace (1) or not (0). This state space has some flaws. For instance, the state $(21, 1, S_d)$ may or may not be a blackjack (depending on how many cards have been drawn). Moreover, in the finite deck setting, this representation does not keep track of the remaining number of cards in the deck.

In the numerical experiments, both \mathcal{S}_1 and \mathcal{S}_2 will be considered.

3 Numerical experiments for blackjack

A Python code for reinforcement learning of blackjack is provided. Before you begin, you need to do the following:

1. Install *Gymnasium* to your Python environment, see <https://gymnasium.farama.org/>, e.g. using `pip` in your terminal: `pip install gymnasium`.
2. Download the project Python files, available on the Canvas page under **Files**.
3. In your code directory, create subdirectories named `data` and `figures`. The output of the learning algorithms will be saved in these directories.

The main program is contained in `run_RL.py`, which is set up to train the agent for 10^n episodes ($n = 3$ as the default, increase this for your report to adequately assess the long-term performance for all algorithms) on a specified number of decks, 1, 2, 6, 8, and infinitely many decks. The larger state space, \mathcal{S}_1 , called the *extended state space*, is encoded in the environment `BlackjackEnvExtend`, whereas the smaller state space \mathcal{S}_2 , called the *sum-state space*, is encoded in the environment `BlackjackEnvBase`.

The learning algorithms are partially implemented in the file `RL.py` in the functions `learn_MC` for *on-policy Monte Carlo learning* and `learn_Q` for *Q-learning*. In both of these programs, you need to add the suitable code for learning the action-values for each state-action pair.

Assignments

- (a) *On-policy Monte Carlo learning.* At the indicated place in the function `learn_MC`, add the code for updating the reward in the variable `avg_reward`, and for updating the action-values in the dictionary named `Q`.
- (b) *Q-learning.* At the indicated place in the function `learn_Q`, add the code for computing the learning rate and for updating the action-values in the dictionary named `Q`. Optional challenge: Try also implementing the *SARSA* or *Expected SARSA* algorithm here.
- (c) For the following three cases:
 - (1) on-policy Monte Carlo learning on the state space \mathcal{S}_1 ,
 - (2) Q-learning on the state space \mathcal{S}_1 ,
 - (3) Q-learning on the state space \mathcal{S}_2 ,

compare the performance of the learning algorithms. Which one achieves better learning? For (2) and (3), experiment on the learning rate, to find a suitable decay rate of the form $\alpha_n = c n^{-\omega}$, $0 \leq \omega < 1$. Also, explore the impact of having a decaying exploration rate ϵ . The exploration rate is the probability of not choosing the optimal policy. Optional challenge: Does the current evaluation give a fair assessment of the target policy discovered by Q-learning? Can you improve it?

- (d) For the best learning algorithm, provide a close to optimal strategy by specifying a table of actions (“hit” or “stay”), depending on the agent’s card sum, the dealer’s card sum, and whether there is a usable ace or not.

Report

Your group must hand in a report, by uploading it on Canvas, containing the following items:

- A mathematical description of the on-policy Monte Carlo learning algorithm on which the solution in (a) is based.
- A mathematical description of the Q-learning algorithm on which the solution in (b) is based.
- Plots comparing the performance of the algorithms as mentioned in (c), as well as relevant plots that support your evaluation of the learning rate.
- A description of a close to optimal strategy by specifying a table of actions as mentioned in (d).